



1 SUMMARY

Let A be an $n \times n$ matrix with a symmetric sparsity pattern. HSL_MC70 computes a **nested dissection ordering** of A that is suitable for use with a sparse direct solver. The algorithm allows for some dense or nearly dense rows and columns in A .

The algorithm partitions the rows/columns of A into 3 sets such that reordering the rows/columns to respect to their partitions yields a symmetric matrix of the form

$$\left(\begin{array}{c|c|c} A_1 & 0 & S_1^T \\ \hline 0 & A_2 & S_2^T \\ \hline S_1 & S_2 & S \end{array} \right). \quad (1.1)$$

The rows of S are given an arbitrary order. If the dimension of A_1 (A_2) is smaller than some predefined value, the rows of A_1 (A_2) will be ordered using an approximate minimum degree algorithm; otherwise, the rows/columns of A_1 (A_2) will be partitioned to form another matrix with the above structure and the algorithm will continue to be applied in a recursive manner.

ATTRIBUTES — **Version:** 1.0.0. **Types:** Real (single, double). **Uses:** HSL_MC78 and HSL_MC79. **Language:** Fortran 2003 subset (F95+TR155581) **Date:** January 2014. **Origin:** I.S. Duff, J.A. Scott and H.S. Thorne, Rutherford Appleton Laboratory.

2 HOW TO USE THE PACKAGE

2.1 Calling sequences

Access to the package requires a USE statement of the form

```
USE HSL_MC70_integer
```

2.2 The derived data types

For each problem, the user must employ the derived types defined by the module to declare scalars of the types MC70_control and MC70_info. The following pseudocode illustrates this.

```
use HSL_MC70_integer
...
type (MC70_control) :: control
type (MC70_info) :: info
...
```

The components of MC70_control and MC70_info are explained in Sections 3.3 and 3.4.

3 THE ARGUMENT LISTS

3.1 Input of the matrix A

The user must supply the pattern of either the **lower triangular part** of the matrix A or the **lower and upper triangular part** of A in a compressed sparse column format. There is no requirement that zero entries on the diagonal are

explicitly included. **No checks** are made on the user's data. It is important to note that any out-of-range entries or duplicates may cause HSL_MC70 to fail in an unpredictable way. Before using HSL_MC70, the HSL package HSL_MC69 may be used to check for errors and to handle duplicates (HSL_MC69 sums them) and out-of-range entries (HSL_MC69 removes them).

If the user's data is held using another standard sparse matrix format (such as coordinate format or sparse compressed row format), we recommend using a conversion routine from HSL_MC69 to put the data into the required format. The input of A is illustrated in Section 6.

3.2 To compute a nested dissection ordering

If the user has the **lower triangular part of A** held in compressed sparse columns format, a call of the following form should be made:

```
CALL mc70_order(n,ptr,row,perm,control,info)
```

If the user has the **lower and upper triangular parts of A** held in compressed sparse columns format, a call of the following form should be made:

```
CALL mc70_order_full(n,ptr,row,perm,control,info)
```

n is an INTEGER scalar with INTENT(IN). On entry it must hold the order n of A . **Restriction:** $n \geq 1$.

ptr is an INTEGER array of rank one with INTENT(IN) and size $n+1$. It must be set by the user so that $ptr(j)$ is the position in row of the first entry in column j ($j=1,2,\dots,n$) and $ptr(n+1)$ must be set to one more than the total number of entries.

row is an INTEGER array of rank one with INTENT(IN) and size $ptr(n+1)-1$. On a call to `mc70_order`, it must be set by the user so that $row(1:ptr(n+1)-1)$ holds the row indices of the entries in the **lower triangular part** of A ; on a call to `mc70_order_full`, it must be set by the user so that $row(1:ptr(n+1)-1)$ holds the row indices of the entries in the **lower and upper triangular parts** of A . The entries in a single column must be contiguous. The entries of column j must precede those of column $j+1$ ($j=1,2,\dots,n-1$), and there must be no wasted space between the columns. Row indices within a column may be in any order. Diagonal entries are ignored.

$perm$ is an INTEGER array with INTENT(OUT) and size n . On exit, $perm$ holds the nested dissection ordering. The position of variable i in the nested dissection ordering is $perm(i)$, $i=1,2,\dots,n$.

$control$ is a scalar of type `MC70_control` with INTENT(IN). Its components control the action, as explained in Section 3.3.

$info$ is a scalar of type `MC70_info` with INTENT(OUT). Its components hold information, as explained in Section 3.4.

3.3 The control derived data type

The derived data type `MC70_control` is used to control the action. The user must declare a structure of type `MC70_control`. The components, which are automatically given default values in the definition of the type, are:

Printing controls

`print_level` is a scalar of type INTEGER that is used to controls the level of printing. The different levels are:

- < 0 No printing.
- $= 0$ Error messages only.

- = 1 As 0, plus basic diagnostic printing.
- > 1 As 1, plus some additional diagnostic printing.

The default is `print_level=0`.

`unit_diagnostics` is a scalar of type `INTEGER` that holds the unit number for diagnostic printing. Printing is suppressed if `unit_diagnostics < 0`. The default is `unit_diagnostics=6`.

`unit_error` is a scalar of type `INTEGER` that holds the unit number for error messages. Printing of error messages is suppressed if `unit_error < 0`. The default is `unit_error=6`.

Other controls

`max_levels` is an `INTEGER` scalar that holds the maximum number of nested dissection levels. If the maximum number of nested dissection levels is reached, the corresponding matrix in the nested dissection hierarchy will be ordered using an approximate minimum degree algorithm. If `max_levels = 0`, the ordering of the input matrix is computed using an approximate minimum degree algorithm. Values less than 0 are treated as 0. The default is `max_levels=20`.

`nd_switch` is an `INTEGER` scalar that is used in the criteria for determining what happens a matrix within the nested dissection hierarchy. If the matrix order is greater than `nd_switch` and the maximum number of nested dissection levels has not been reached, the matrix will be partitioned; otherwise, it will be ordered using an approximated minimum degree algorithm. Values less than 2 are treated as 2. The default is `nd_switch=50`.

`ratio` is a `DOUBLE PRECISION` scalar. The partitioning and refinement methods aim to find partitions (1.1) such that $\max(n_1, n_2) < \min(n_1, n_2) \times \text{ratio}$, where A_1 has order n_1 and A_2 has order n_2 . If several candidate partitions satisfy this requirement, then the partition with the smallest value of $\frac{n_s}{n_1 n_2}$ is chosen, where S has order n_s ; if none of the candidate partitions satisfy this requirement, the partition that has the smallest value of $\frac{n_s}{n_1 n_2}$ is chosen. Decreasing `ratio` will, in general, result in a nested dissection ordering that is more amenable to parallel direct solvers; increasing `ratio` will, in general, reduce the number of non-zeros in the Cholesky factorization of the reordered matrix. The default is `ratio=1.5`.

`remove_dense` is a `LOGICAL` scalar. If `remove_dense = .true.`, then the input matrix is searched for dense (or nearly dense) rows and columns, and the nested dissection algorithm is applied to the matrix that results when these rows and columns are removed. Dense rows/columns are placed at the end of the ordering. If `remove_dense = .false.`, then the input matrix is not searched for dense rows. The default is `remove_dense=.true.`

3.4 The derived data type for holding information

The derived data type `MC70_info` is used to hold parameters that give information about the algorithm. The components of `MC70_info` (in alphabetical order) are:

`flag` is a scalar of type `INTEGER` that gives the exit status of the algorithm (details in Section 3.5).

`ndense` is a scalar of type `INTEGER` that holds the number of rows/columns in the input matrix that were determined to be dense. If `control%remove_dense = .false.`, the input matrix is not checked for dense rows.

`stat` is a scalar of type `INTEGER` that holds the Fortran `stat` parameter.

3.5 Error diagnostics

On successful completion, MC70_nested will exit with `info%flag` set to 0. Other values for `info%flag` are associated with a fatal error. Possible values are:

- 1 memory allocation failed.
- 2 memory deallocation failed.
- 3 $n \leq 0$.

4 GENERAL INFORMATION

Workspace: Provided automatically by the module.

Other modules used directly: HSL_MC78 and HSL_MC79.

Input/output: Error and diagnostic messages. Error messages on unit `control%unit_err` and diagnostic messages on unit `control%unit_diagnostics`. These have default value 6; printing of these messages is suppressed if the relevant unit number is negative or if `control%print_level` is negative.

Restrictions: $n \geq 1$.

Portability: Fortran 2003 subset (F95+TR155581)

5 METHOD

Given a symmetric matrix A of order n , HSL_MC70 preprocesses the matrix to (optionally) remove dense or almost dense rows/columns, returning \bar{A} . The matrix \bar{A} is then compressed using HSL_MC78 to give a symmetric matrix \tilde{A} with order \tilde{n} . This matrix is passed to the recursive nested dissection algorithm, Algorithm 1, and the resulting permutation matrix is converted into the corresponding elimination ordering. Having formed a nested dissection ordering for \tilde{A} , the ordering is mapped to give an ordering for \bar{A} . The dense rows are appended to the end of the ordering.

Algorithm 1 Nested dissection algorithm

recursive subroutine nested_dissection(A, P)

Input: symmetric matrix A of order n

Output: permutation matrix P

if $n < \text{nd_switch}$ **then**

 Form the AMD elimination ordering, p , for A and return its equivalent permutation matrix

else

 Partition the matrix: compute P , a permutation matrix, such that $P^T A P$ has the form (1.1)

 Call **nested_dissection**(A_1, P_1)

 Call **nested_dissection**(A_2, P_2)

 Perform the update $P = P Q$, where

$$Q = \begin{pmatrix} P_1 & & \\ & P_2 & \\ & & I \end{pmatrix}$$

end if

5.1 Partitioning a matrix**5.2 Refinement****6 EXAMPLE OF USE**